

Project 1: Realizing Concurrency using Linux Processes and Threads

Qinya Li

Department of Computer Science and Engineering

Shanghai Jiao Tong University

Spring 2023

Objective

- Design and develop systems programs using C/C++
- Effectively use Linux system calls for process control and management, especially, *fork*, *exec*, *wait*, *pipe* and *kill* system call API.
- Concurrent execution of processes.
- Use Posix Pthread library for concurrency.
- Use graphic tools or program to analyze data.

Project 1

- Copy
- Shell
- Matrix multiplication

Project 1: Copy

- Copy file from one to another

Project 1: Copy

- Copy file from one to another
- Use two processes
- Processes communicate using pipe system call

Project 1: Copy

- Copy file from one to another
- Use two processes
- Processes communicate using pipe system call
- Use buffer in order to copy faster

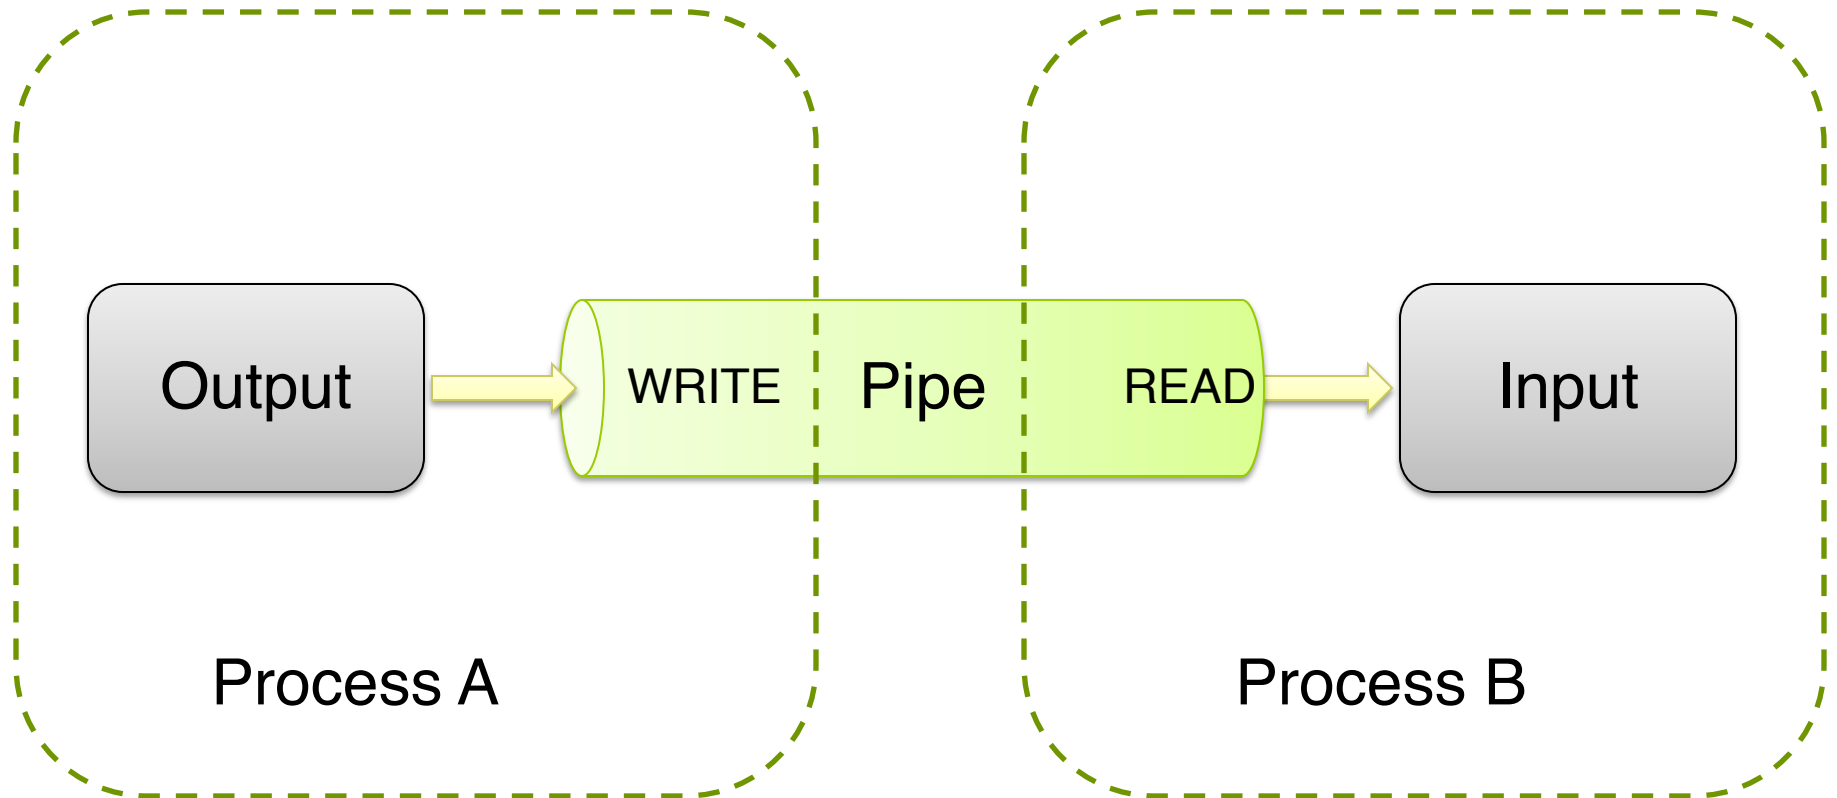
Project 1: Copy

- Copy file from one to another
- Use two processes
- Processes communicate using pipe system call
- Use buffer in order to copy faster
- Use system calls to get executing time

Project 1: Copy

- Copy file from one to another
- Use two processes
- Processes communicate using pipe system call
- Use buffer in order to copy faster
- Use system calls to get executing time
- Use graphic tools or programs to analyze the result

Ordinary Pipe



Project 1: Copy Tips

- Open a file for reading and check for errors:

```
FILE *src;
```

```
// Open source file
```

```
src = fopen (argv[1], "r");
```

```
// Check for file error
```

```
if (src == NULL) {
```

```
    printf ("Error: Could not open file '%s'.\n", argv[1]);
```

```
    exit(-1);
```

```
}
```

- “r”: Opens a source file for reading

Project 1 Write Tips

- Open a file for writing and check for errors:

```
FILE *dest;
```

```
// Open destination file
```

```
dest = fopen(argv[2], "w+");
```

```
// Check for file error
```

```
if (dest == NULL) {  
    printf("Error: Could not open file '%s'.\n", argv[2]);  
    fclose(src);  
    exit(-1);  
}
```

- “w+”: Opens an empty file for both reading and writing. If the given file exists, its contents are destroyed.

Project 1 Tips

■ Read and write to a file

- Read/Write a block of chars:

size_t fread(void **buffer*, size_t *size*, size_t *count*, FILE **stream*);

- ▶ **fread** returns the number of full items actually read

size_t fwrite(const void **buffer*, size_t *size*, size_t *count*, FILE **stream*);

- ▶ **fwrite** returns the number of full items actually written

- Read an individual char:

int fgetc(FILE **stream*);

int fputc(int *c*, FILE **stream*);

Project 1 Tips

- Don't forget to check for EOF (with `fgetc`) or size of buffer written (with `fread/fwrite`)
- Don't forget to close your file with `fclose`

Project 1 Problem 1 Tips

■ Timing I

```
clock_t start, end;  
double elapsed;  
start = clock();  
.....  
end = clock();  
elapsed = ((double) (end - start)) /  
           CLOCKS_PER_SEC * 1000;  
printf("Time used: %f millisecond\n", elapsed);
```

Project 1 Problem 1 Tips

■ Timing II

```
struct timeval startTime, endTime;
struct timezone tz;
struct tm *tm;
gettimeofday(&startTime, &tz);
.....
gettimeofday(&endTime, &tz);
long run_time_in_microseconds;
run_time_in_microseconds = endTime.tv_usec - startTime.tv_usec;
printf("Executable time for quick sort single threaded is %d
microseconds.\n",run_time_in_microseconds);
```

```
struct timeval{
    long int tv_sec; // seconds
    long int tv_usec; // microseconds
}
```

Using fork()

■ Fork the child process:

```
pid_t ForkPID;  
ForkPID = fork();
```

■ Write a quick switch statement:

```
switch (ForkPID) {  
    // -1, fork failure  
    case -1:  
        printf("Error: Failed to fork.\n"); break;  
    // 0, this is the child process  
    case 0:  
        .....  
        break;  
    // > 0, parent process and the PID is the child's PID  
    default:  
        .....  
}
```


Using Pipe – Part 1

- First, create a pipe and check for errors

```
int mypipe[2];  
if (pipe(mypipe)) {  
    fprintf (stderr, "Pipe failed.\n");  
    return -1;  
}
```

- Second, fork your processes
- Third, close the pipes you don't need in that process
 - reader should close(mypipe[1]);
 - writer should close(mypipe[0]);

Using Pipe – Part 2

- Fourth, the writer should write the data to the pipe
 - `ssize_t write(int pipe_id, const void *buf, size_t count)`
 - `write(mypipe[1],&c,1);`
- Fifth, the reader reads from the data from the pipe:
 - `ssize_t read(int pipe_id, const void *buf, size_t count)`
 - `while (read(mypipe[0],&c,1)>0) {`
`//do something, loop will exit when WRITER closes pipe }`
- Sixth, when writer is done with the pipe, close it
 - `close(mypipe[1]); //EOF is sent to reader`
- Fifth, when reader receives EOF from closed pipe, close the pipe and exit your polling loop
 - `close(mypipe[0]); //all pipes should be closed now`

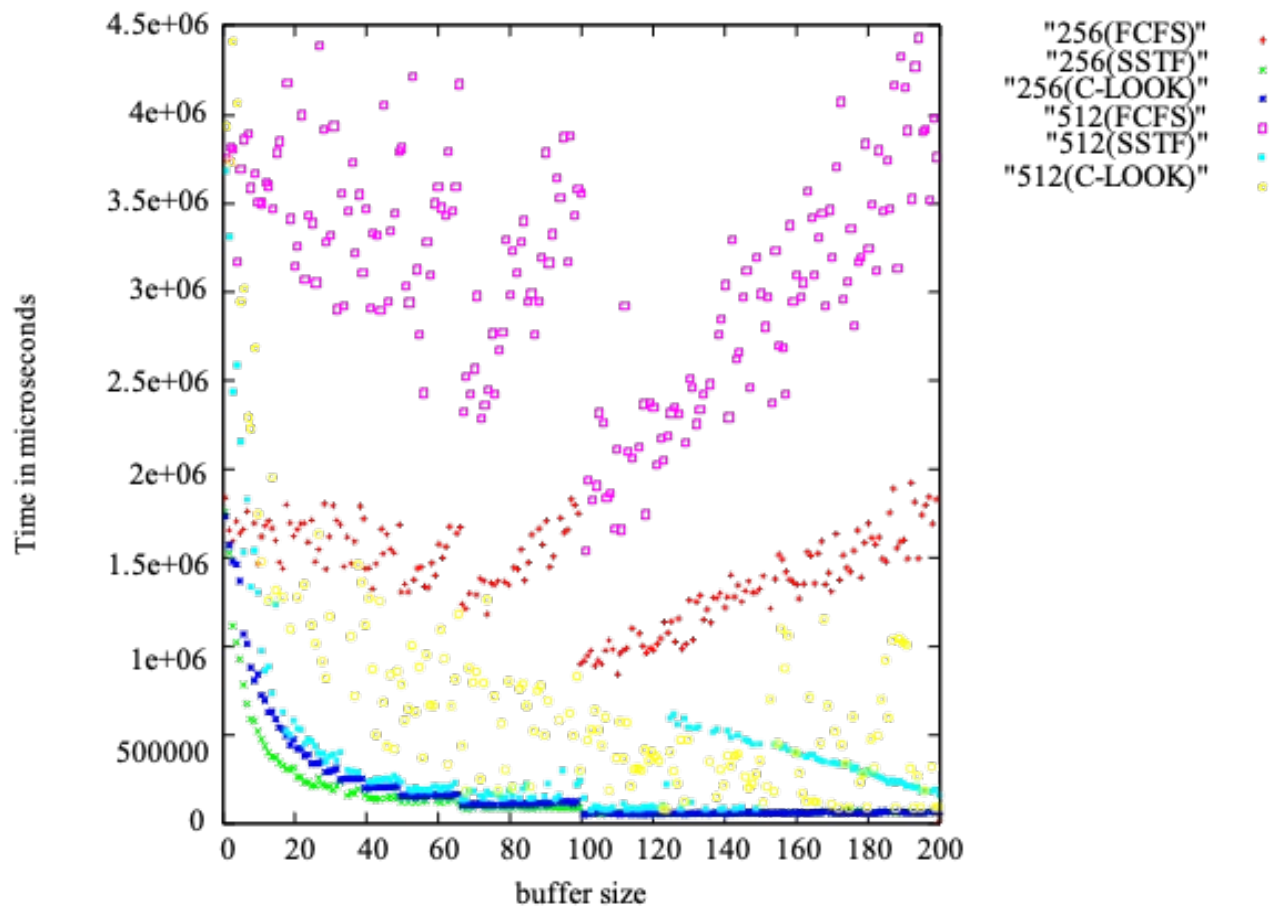
Example

```
lqy@ubuntu: ~/OSProject
File Edit View Search Terminal Help
lqy@ubuntu:~/OSProject$ gcc copy.c -o copy
lqy@ubuntu:~/OSProject$ ls
copy  copy.c  copy.out  dest.txt  sample_data.in  src.txt
lqy@ubuntu:~/OSProject$
```

```
lqy@ubuntu: ~/OSProject
File Edit View Search Terminal Help
lqy@ubuntu:~/OSProject$ ./copy src.txt dest.txt 100
Read file end.
Write file end.
Time used (PipeCopy): 0.384000 millisecond.
lqy@ubuntu:~/OSProject$ ./copy src.txt dest.txt 1000
Read file end.
Write file end.
Time used (PipeCopy): 0.248000 millisecond.
lqy@ubuntu:~/OSProject$
```

Gnuplot

- Use graphic tools or program to analyze data
- <http://www.gnuplot.info/>



Project 1: Shell

- Write a shell-like program as a server

Project 1: Shell

- Write a shell-like program as a server
- Handle the commands with arguments and the commands connected by pipes

Project 1: Shell

- Write a shell-like program as a server
- Handle the commands with arguments and the commands connected by pipes
- Try to support more than one client

Project 1: Shell

- Write a shell-like program as a server
- Handle the commands with arguments and the commands connected by pipes
- Try to support more than one client
- **Do NOT** use the linux shell command line processor to run your command line, you should handle the arguments by yourself

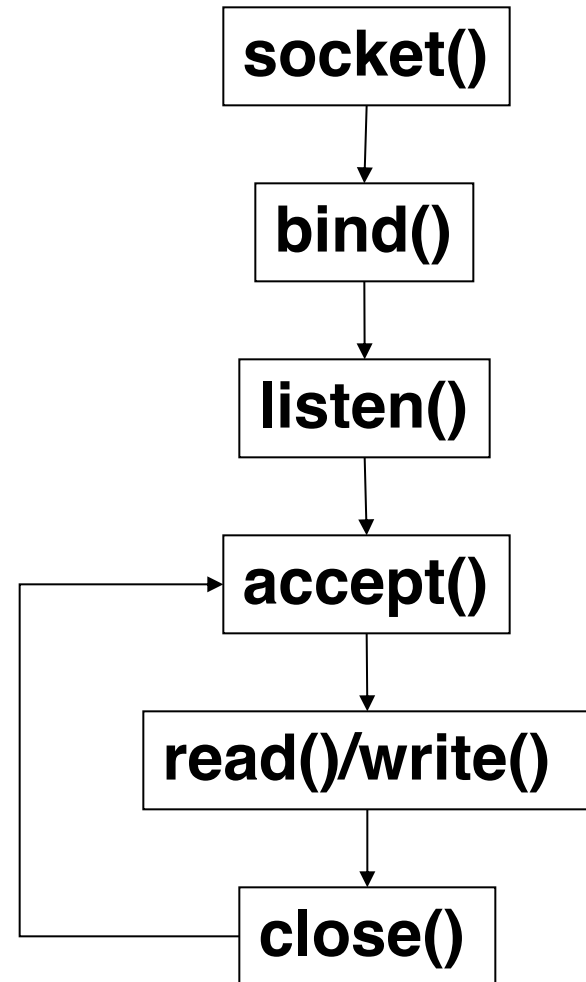
Resources

- Here are a few good pages that can help:
 - http://www.linuxhowtos.org/C_C++/socket.htm
 - http://en.wikipedia.org/wiki/Internet_socket
 - http://en.wikipedia.org/wiki/Berkeley_sockets
- Sample client/server codes available on course webpage
- You can use this code on these pages, as you will need to change them to suit your own purposes

Socket Programming

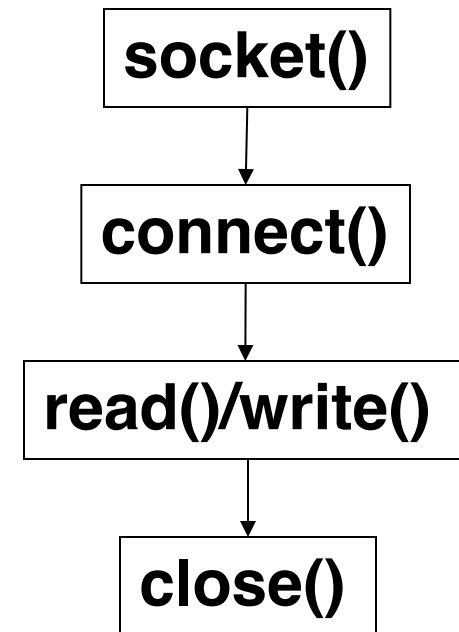
Steps to Create Server Side

1. Create a socket with the **socket()** system call
2. Bind the socket to an address using the **bind()** system call.
3. Listen for connections with the **listen()** system call
4. Accept a connection with the **accept()** system call (This call typically blocks until a client connects with the server)
5. Send and receive data with **read()** and **write()** system calls
6. Close connection with **close()** system call

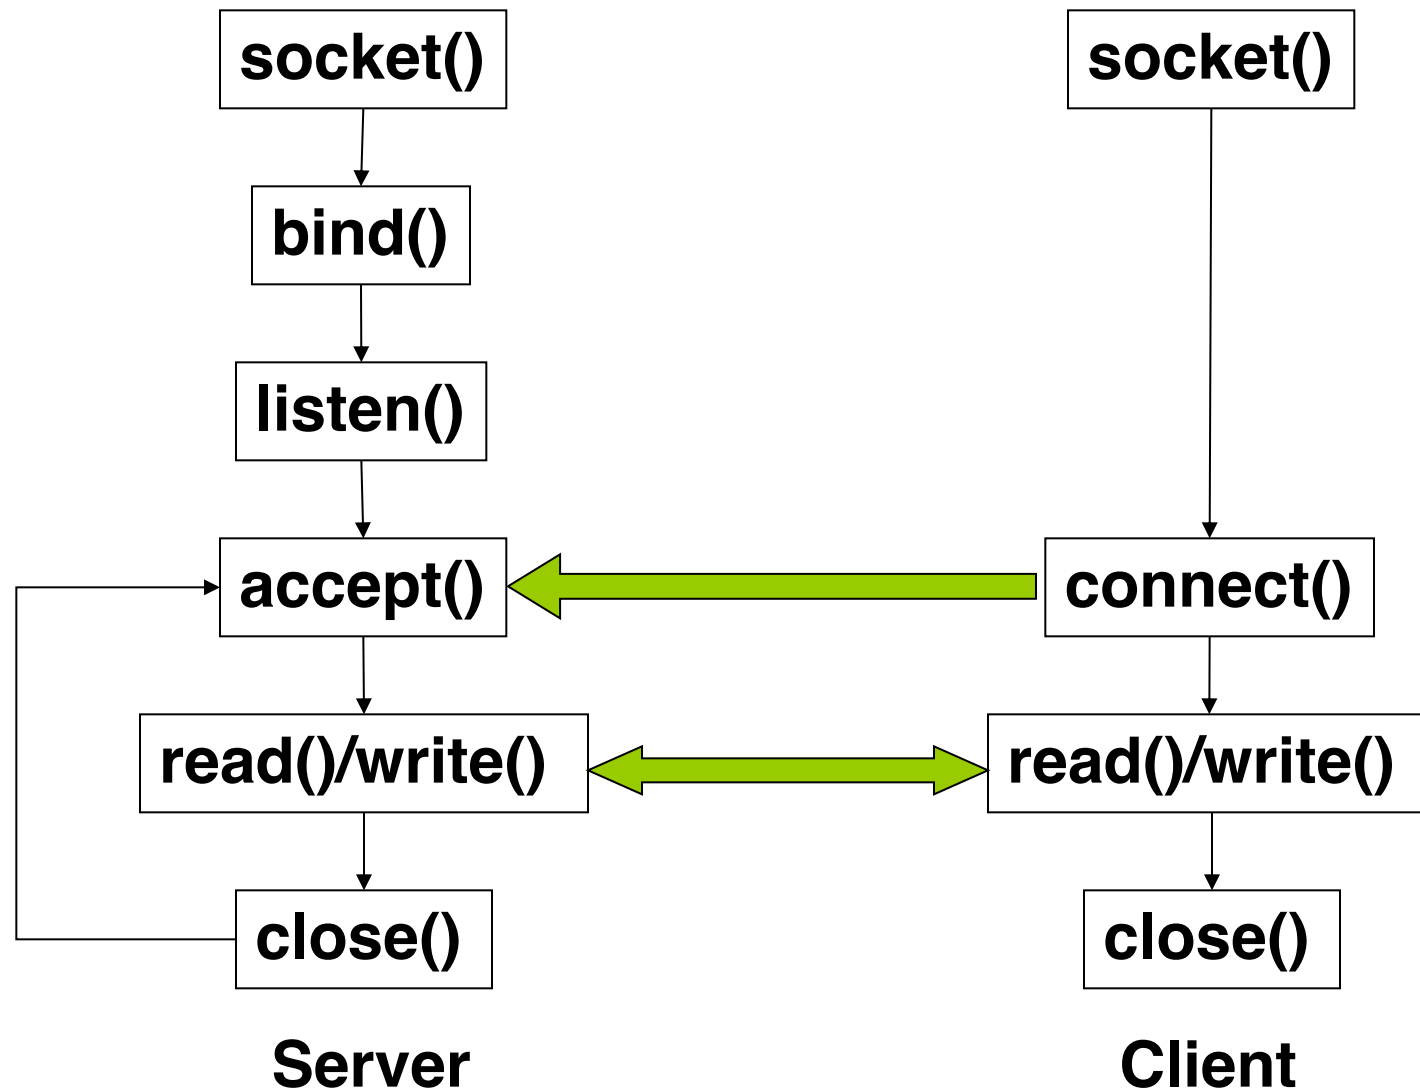


Steps to Create Client Side

1. Create a socket with the **socket()** system call
2. Connect the socket to the address of the server using the **connect()** system call
3. Send and receive data with **read()** and **write()** system calls.
4. Close the socket with **close()** system call



Interaction Between Client and Server



Internet Domain Socket

■ IP address:

- 32 bits (IPv4) or 128 bits (IPv6)
- C/S work on same host: just use **localhost**

■ Port

- 16 bit unsigned integer
- Lower numbers are reserved for standard services

■ Transport layer protocol: TCP / UDP

Headers

- `#include <stdio.h>`
- `#include <stdlib.h>`
- `#include <string.h>`
- `#include <sstream>`
- `#include <unistd.h>`
- `#include <sys/types.h>`
 - Definitions of a number of data types used in system calls
- `#include <sys/socket.h>`
 - Definitions of structures needed for sockets
- `#include <netinet/in.h>`
 - Constants and structures needed for internet domain addresses

Creating Socket

```
int sockfd
sockfd = socket(AF_INET, SOCK_STREAM, 0);
if (sockfd < 0) {
    perror("ERROR opening socket");
    exit(2);
}
```

- AF_INET: address domain
- SOCK_STREAM: stream socket, characters are read in a continuous stream as if from a file or pipe
- 0: protocol. The operating system chooses the most appropriate protocol. It will choose TCP for stream sockets.

Binding Socket

```
struct sockaddr_in serv_addr;  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_addr.s_addr = htonl(INADDR_ANY);  
serv_addr.sin_port = htons(BASIC_SERVER_PORT);  
bind(sockfd, (sockaddr*) &serv_addr, sizeof(serv_addr));  
//error check
```

- INADDR_ANY: get IP address of the host automatically
- htonl, htons: data format conversion
- bind(): binds a socket to an address

Listening and Accepting Connection

```
listen(sockfd, 5);
```

- **listen()**: allows the server to listen on the socket for connections, with a backlog queue of size 5.

```
int client_sockfd;  
struct sockaddr_in client_addr;  
int len = sizeof(client_addr);  
client_sockfd = accept(sockfd, (sockaddr *) &client_addr,  
&len);  
    //error check
```

- **accept()**: block process until a client connects to the server. It returns a new socket file descriptor, if the connection is created.

Reading and Writing

```
char buf[1024];  
int nread = read(client_sockfd, buf, 1024);
```

read(): reads from the socket

```
write(client_sockfd, buf, len);
```

write(): writes to the socket

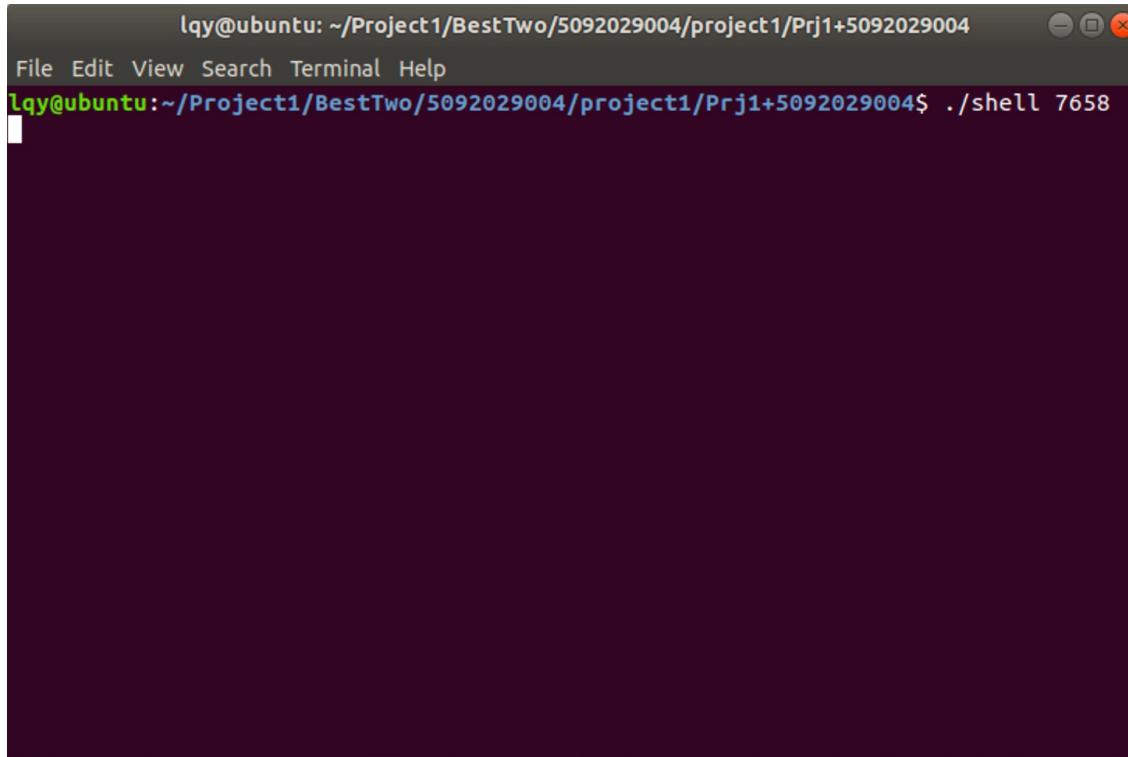
```
close(client_sockfd);
```

close(): closes the socket

Connecting A Client to A Server

```
int sockfd;  
sockfd = socket(AF_INET, SOCK_STREAM, 0);  
    // error check  
struct sockaddr_in serv_addr;  
struct hostent *host;  
serv_addr.sin_family = AF_INET;  
host = gethostbyname(argv[1]);  
    // error check  
memcpy(&serv_addr.sin_addr.s_addr, host->h_addr,  
    host->h_length);  
serv_addr.sin_port = htons(BASIC_SERVER_PORT);  
connect(sockfd, (sockaddr *) &serv_addr, sizeof(serv_addr))  
    // error check
```

Example



A terminal window with a dark background and light text. The title bar at the top reads "lqy@ubuntu: ~/Project1/BestTwo/5092029004/project1/Prj1+5092029004". Below the title bar is a menu bar with "File", "Edit", "View", "Search", "Terminal", and "Help". The main area of the terminal shows a prompt "lqy@ubuntu:~/Project1/BestTwo/5092029004/project1/Prj1+5092029004\$" followed by the command "./shell 7658". The command has been executed, and the prompt is now on a new line.

```
lqy@ubuntu: ~/Project1/BestTwo/5092029004/project1/Prj1+5092029004
File Edit View Search Terminal Help
lqy@ubuntu:~/Project1/BestTwo/5092029004/project1/Prj1+5092029004$ ./shell 7658
lqy@ubuntu:~/Project1/BestTwo/5092029004/project1/Prj1+5092029004$
```

Example

```
lqy@ubuntu: ~/Project1/BestTwo/5092029004/project1/Prj1+5092029004
File Edit View Search Terminal Help
lqy@ubuntu:~/Project1/BestTwo/5092029004/project1/Prj1+5092029004$
lqy@ubuntu:~/Project1/BestTwo/5092029004/project1/Prj1+5092029004$ ./shell 7658
New child process (seq=1)
```

```
lqy@ubuntu: ~
File Edit View Search Terminal Help
lqy@ubuntu:~$ telnet localhost 7658
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
Welcome to my shell!
/home/lqy/Project1/BestTwo/5092029004/project1/Prj1+5092029004$
```

Example

```
lqy@ubuntu: ~/Project1/BestTwo/5092029004/project1/Prj1+5092029004
File Edit View Search Terminal Help
lqy@ubuntu:~/Project1/BestTwo/5092029004/project1/Prj1+5092029004$ ./shell 7658
New child process (seq=1)
New child process (seq=2)
```

```
lqy@ubuntu: ~
File Edit View Search Terminal Help
lqy@ubuntu:~$ telnet localhost 7658
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
Welcome to my shell!
/home/lqy/Project1/BestTwo/5092029004/project1/Prj1+5092029004$
```

Project 1 Shell Tips – Part 1

- Shell must specify an input loop as described, so you should implement a prompt-input-execute loop
- Sample test run:

```
lqy@ubuntu: ~  
File Edit View Search Terminal Help  
src.txt  
test  
/home/lqy/Project1/BestTwo/5092029004/project1/Prj1+5092029004$ pwd  
/home/lqy/Project1/BestTwo/5092029004/project1/Prj1+5092029004  
/home/lqy/Project1/BestTwo/5092029004/project1/Prj1+5092029004$ ls -l  
total 252  
-rw-rw-r-- 1 lqy lqy 2152 Jun 19 2012 copy.c  
-rw-rw-r-- 1 lqy lqy 15770 Jun 19 2012 copy.ods  
-rwxrwxr-x 1 lqy lqy 13008 Feb 16 22:42 copy.out  
-rw-rw-r-- 1 lqy lqy 86 Feb 16 22:44 des.txt  
-rw-rw-r-- 1 lqy lqy 460 Jun 19 2012 Makefile  
-rwxrwxr-x 1 lqy lqy 12313 Jun 19 2012 multi  
-rwxrwxr-x 1 lqy lqy 314 Jun 19 2012 multi-benchmark.sh  
-rw-rw-r-- 1 lqy lqy 5988 Jun 19 2012 multi.c  
-rw-rw-r-- 1 lqy lqy 38612 Jun 19 2012 multi.ods  
-rw-rw-r-- 1 lqy lqy 83673 Jun 19 2012 Performance Analysis.pdf  
-rw-rw-r-- 1 lqy lqy 476 Jun 19 2012 Prj1README  
-rwxrwxr-x 1 lqy lqy 13712 Feb 19 19:31 shell  
-rw-rw-r-- 1 lqy lqy 6236 Jun 19 2012 shell.c  
-rwxrwxr-x 1 lqy lqy 13712 Feb 17 02:18 shell.out  
-rw-rw-r-- 1 lqy lqy 3828 Jun 19 2012 single.c  
-rw-rw-r-- 1 lqy lqy 86 Feb 16 22:43 src.txt  
-rw-rw-r-- 1 lqy lqy 1772 Jun 19 2012 test  
/home/lqy/Project1/BestTwo/5092029004/project1/Prj1+5092029004$
```


Project 1 Shell Tips – Part 2

- Fork a child process for executing each command (the exec system call will overlay the child code and use the child process space)
- Sample code run in the child process:

```
if (execvp(command_array[0],command_array) == -1) {  
    printf("Error: running command: '%s'\n", line);  
    exit(0);}
```

For example:

```
line = "ls -l"
```

```
command_array[0] = "ls"
```

```
command_array[1] = "-l"
```

```
command_array[2] = NULL
```

Project 1 Shell Tips – Part 3

■ Parsing command

```
int parseLine(char *line, char *command_array[]) {  
    char *p;  
    int count=0;  
    p = strtok(line, " ");  
    while (p && strcmp(p,"")!=0) {  
        command_array[count] = p;  
        count++;  
        p = strtok(NULL, " ");  
    }  
    return count;  
}
```

Project 1 Shell Tips – Part 4

- Very IMPORTANT: you do not need to handle “*” or “?”.
- Handling “exit” command!

Project 1 Shell Tips – Part 5

- In the parent process, wait for a child process to complete (exit) before prompting again:

```
// wait (or block current process) until a forked child finished  
int r;  
wait(&r);
```

An alternative way (or see **waitpid()** call):

```
// do not exit until ALL child processes are finished  
// be aware, this may behave differently depending on the  
// machine you use or how fast things are  
while (wait(&status) != 0) {}
```

Project 1 Shell Tips – Part 6

- Needs to add:
 - Processing the pipe “|” symbol in the command line
 - Piping/redirecting data from one process to another
- If pipe without command after it?
 - throw an error message, continue loop
- If pipe at front of command line?
 - throw an error message, continue loop

Project 1 Shell Tips – Part 7

- To “redirect” input/output, use dup command
 - We prefer to use **dup2()** command
 - `dup2(int oldfd, int newfd);`
- Steps for programming:
 - First, have main program process both sides of the pipe into a command array
 - ▶ create a pipe (just like in problem 1)
 - ▶ fork a child process A
 - ▶ fork a child process B
 - ▶ close both ends of the pipe in the parent

Project 1 Shell Tips – Part 8

■ Steps for programming, cont.:

● Second, in child process A

- ▶ close the READ pipe, and standard output
 - **close(mypipe[0]);**
 - **close(1);**
- ▶ dup2 the WRITE pipe onto standard output AND THEN close the WRITE pipe
 - **dup2(mypipe[1],STDOUT_FILENO);**
 - **close(mypipe[1]);**
- ▶ exec the command after pipe(with arguments)
 - if (execvp(command_arrayB[0],command_arrayB) == -1)**
 - { printf("Error: running command: '%s'\n", line); }**

Project 1 Shell Tips – Part 9

■ Steps for programming, cont.:

● Third, in child process B

- ▶ close the WRITE pipe, and standard input
 - `close(mypipe[1]);`
 - `close(0);`
- ▶ dup2 the READ pipe onto standard input AND THEN close the READ pipe
 - `dup2(mypipe[0],STDIN_FILENO);`
 - `close(mypipe[0]);`
- ▶ exec the command before pipe (with arguments)
 - if (`execvp(command_arrayA[0],command_arrayA) == -1`)
 - { `printf("Error: running command: '%s'\n",line);` }

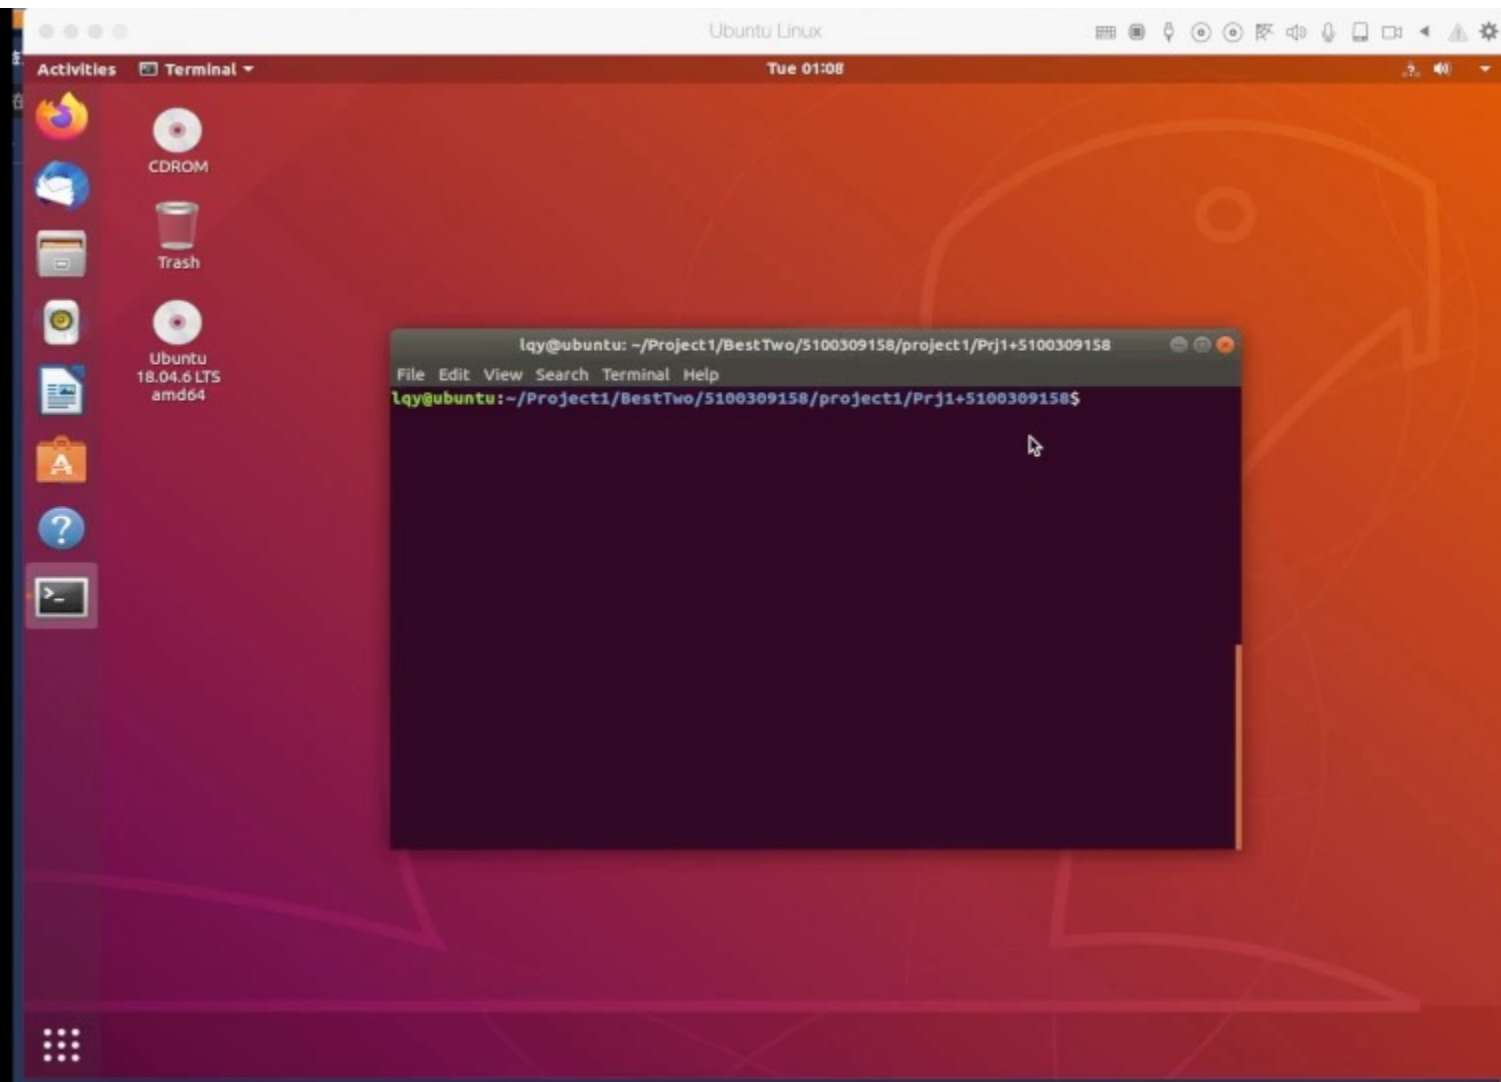
Project 1 Shell Tips – Part 10

- Steps for programming, cont.:
 - Finally, have the parent wait until child processes are all done.
 - ▶ **wait(&r);** // hold back new prompt until old command finishes

Project 1 Shell NoNo's

- Do NOT use the linux shell command line processor to run your command line. You will get zero points for that solution.
- For example, do not do the following to run a “ls -l *.c” command:
 - `execl("/usr/bin/sh", "sh", "-c", "ls -l *.c", NULL);`
 - This is using the shell “sh” to do all the heavy work and thus is not acceptable

Sample Shell



Project 1: Matrix multiplication

- Find a algorithm of matrix multiplication

Project 1: Matrix multiplication

- Find a algorithm of matrix multiplication
- Implement a single-thread program and a multi-thread program

Project 1: Matrix multiplication

- Find a algorithm of matrix multiplication
- Implement a single-thread program and a multi-thread program
- Compare the performance of different implementations

Project 1 Pthread 1

- Header needed:

```
#include <pthread.h>
```

- Preparation:

```
// create a thread 't1'
```

```
pthread_t t1;
```

```
pthread_attr_t attr1;
```

```
pthread_attr_init(&attr1);
```

Project 1 Pthread 1

- Header needed:

```
#include <pthread.h>
```

- Preparation:

```
// create a thread 't1'
```

```
pthread_t t1;
```

```
pthread_attr_t attr1;
```

```
pthread_attr_init(&attr1);
```

```
pthread_attr_setdetachstate(&attr1,  
    THREAD_CREATE_JOINABLE);
```

```
THREAD_CREATE_JOINABLE  
THREAD_CREATE_DETACHED
```

```
typedef struct  
{  
    int __detachstate;  
    int __schedpolicy;  
    struct sched_param __schedparam;  
    int __inheritsched;  
    int __scope;  
    size_t __guardsize;  
    int __stackaddr_set;  
    void* __stackaddr;  
    size_t __stacksize;  
} pthread_attr_t;
```


Project 1 Pthread 1

- Header needed:

```
#include <pthread.h>
```

- Preparation:

```
// create a thread 't1'
```

```
pthread_t t1;
```

```
pthread_attr_t attr1;
```

```
pthread_attr_init(&attr1);
```

```
pthread_attr_setdetachstate(&attr1,  
    THREAD_CREATE_JOINABLE);
```

```
THREAD_CREATE_JOINABLE  
THREAD_CREATE_DETACHED
```

```
typedef struct  
{  
    int __detachstate;  
    int __schedpolicy;  
    struct sched_param __schedparam;  
    int __inheritsched;  
    int __scope;  
    size_t __guardsize;  
    int __stackaddr_set;  
    void* __stackaddr;  
    size_t __stacksize;  
} pthread_attr_t;
```

Project 1 Pthread 1

- Header needed:

```
#include <pthread.h>
```

- Preparation:

```
// create a thread 't1'
```

```
pthread_t t1;
```

```
pthread_attr_t attr1;
```

```
pthread_attr_init(&attr1);
```

```
pthread_attr_setdetachstate(&attr1,  
    THREAD_CREATE_JOINABLE);
```

```
pthread_attr_setscope(&attr1, THREAD_SCOPE_SYSTEM);
```

```
typedef struct  
{  
    int __detachstate;  
    int __schedpolicy;  
    struct sched_param __schedparam;  
    int __inheritsched;  
    int __scope;  
    size_t __guardsize;  
    int __stackaddr_set;  
    void* __stackaddr;  
    size_t __stacksize;  
} pthread_attr_t;
```

Project 1 Pthread 2

- Prepare thread arguments

// create an argument structure (specified before this point)

```
struct thread_arguments my_arguments1;
```

- The structure of argument

```
struct thread_arguments  
{  
    int first_paramter;  
    int second_paramter;  
    int return_value;  
};
```

Project 1 Pthread 3

- Launch a thread

```
rc = pthread_create (&t1, &attr1, my_function, &my_arguments1);  
if (rc) {  
    printf("ERROR; return code from pthread_create(t1) is %d\n", rc);  
    exit(-1);}
```

- Sample function for a thread

```
void *my_function (void *args) {  
    thread_arguments *a;           // These two lines are for  
    a = (thread_arguments *)args;  // coding convenience  
    int a1 = a->first_paramter;  
    int b1 = a->second_paramter;  
    a->return_value = a1 + b1;  
    pthread_exit(NULL); // or you can return values here, but  
                        // you are passing a pointer back  
                        // so be careful – i.e. No local variables  
}
```

Project 1 Pthread 4

- Wait on a thread's completion

// join thread 1 and wait for completion

```
void* status1;
```

```
rc = pthread_join (t1, &status1);
```

```
if (rc)
```

```
{
```

```
    printf("ERROR; return code from pthread_join(t1)  
    is %d\n", rc);
```

```
    exit(-1);
```

```
}
```

Project 1 Tips

- Input data file: data.in
- Output file: data.out
- Random: random.out

Some Tips

- You can use “man” if there is a problem with library functions or system calls

Some Tips

- You can use “man” if there is a problem with library functions or system calls
- Manuals are your friends. Try to use manuals rather than other books

Some Tips

- You can use “man” if there is a problem with library functions or system calls
- Manuals are your friends. Try to use manuals rather than other books
- Try to use LATEX to write your report, you will fall in love with it

Some Tips

- You can use “man” if there is a problem with library functions or system calls
- Manuals are your friends. Try to use manuals rather than other books
- Try to use LATEX to write your report, you will fall in love with it
- Enjoy coding, enjoy project